

## Introductory questions

### **Q1: How would you describe your company in terms of its mission and activities?**

I lead Armundia Group, a tech company specializing in digital services for the banking and insurance industries, particularly in the financial investment sector. Founded in 2007 as a spin-off from a private company, we are a product company developing advanced platforms—not a consultancy—focused on supporting strategic innovation. Our mission is to help clients stay competitive in the middle of rapid economic, social, and technological changes. Our clients include private and retail banks, asset and wealth managers, family offices, custodians, and insurance companies.

At Armundia, we embrace a pioneering and integrated approach to innovation, which is reflected in our name, inspired by concepts of harmony and integration. We maintain a strong R&D unit in L'Aquila, collaborating closely with universities to create pragmatic, cutting-edge solutions. Our platforms provide a holistic view of client wealth, covering not only financial instruments but also alternative assets such as insurance policies, precious items, and art. Technologically, we leverage advanced computing, interactive reporting, machine learning, and large language models to enhance strategic workflows and improve the quality of service for advisors and their clients.

### **Q2: How would you describe your professional profile and background?**

I'm the CEO, co-founder, and shareholder of Armundia Group, and I've been part of the company since day one. I have a technical background with a degree in Computer Science and a strong interest in research. I enjoy designing technical-functional solutions that meet the needs of the banking and insurance sectors we serve. While I'm not a tech addict, I see technology as a powerful enabler of strategic visions and concrete opportunities based on our market experience.

### **Q3: Can you briefly describe the software product your company develops, and the role of AI within it?**

We develop two main types of software products. The first focuses on predictive models trained on highly specialized historical market data to assist advisors and wealth managers in making real-time decisions. Explainability is fundamental here, as advisors must clearly justify investment strategies both to themselves and to their clients.

The second involves generative AI and large language models (LLMs) used as copilots to support conversational business intelligence. These tools help bank

employees quickly generate specific reports and analyses from vast relational databases, saving time and improving decision-making.

Additionally, our AI solutions support regulatory compliance by providing precise, concise explanations of financial products according to European MiFID and IDD standards, and they adapt to international regulations as needed. For example, generative AI answers detailed client questions on product documentation, improving transparency and customer trust.

Overall, our goal is to accelerate workflows, provide high-quality information quickly, and enhance the entire production cycle. Internally, we also use productivity tools to speed up coding, analysis, debugging, and testing processes.

#### **Q4: Have you ever heard of privacy or nutrition labels?**

Yes, we are very familiar with privacy aspects due to regulatory requirements. Since we build solutions for banks—which hold extensive data about their clients—data protection is fundamental. GDPR compliance is embedded in our infrastructure and software design, ensuring transparency and proper governance of personal data.

We know about privacy labels mainly from a regulatory standpoint, especially because our clients (the banks) require us to comply with their standards. Recently, as we expand from B2B to also include B2C, enabling end clients to interact more independently, we are becoming more aware of privacy labels, though not yet in a fully structured way. Usually, banks inform us about the necessary labels and compliance requirements.

Overall, our experience with these labels comes mostly from client mandates, and we understand their importance particularly for collecting and managing sensitive data securely, in line with anti-money laundering and other regulations.

### **Evaluating Comprehensibility / Understandability**

#### **Q5: Is the proposed ethics label easily understandable?**

At first, the large and detailed ethics label was a bit challenging for me, but after reading it carefully, I would say it's about 80% understandable and intuitive. With attention, it becomes clear. I believe these kinds of labels are crucial—they protect an equally important asset, which is the client's highly sensitive data. In our industry, banks have extensive knowledge about their clients, so privacy and adherence to specific rules are absolutely essential.

## Evaluating Appropriateness

### **Q6: Does the proposed label appropriately address Human, Societal, and Environmental drivers?**

I have studied the drivers quite well to better understand the context. I find them very suitable for our field, especially autonomy and transparency for users, non-discrimination, and inclusive participation. Energy consumption, emissions, and sustainability are also important to us because many of our clients have moved their software to the cloud, which involves energy use, processing costs, time, money, and environmental impact. Additionally, ESG investments (Environmental, Social, Governance) have become very popular, partly due to tax benefits. This creates a market opportunity and a focus for us on how to properly handle, present, and label these investments in compliance with regulations.

## Evaluating Usefulness

### **Q7: From a developer's perspective, is the proposed ethics label useful for expressing the ethical quality of your product?**

#### **Q7.1: Why?**

Yes, definitely. Here, I'm putting on my technical hat because, for us, the developer isn't just a technician in the generic sense—they're also a creator, a creative person who builds value and tools that concretely impact people's lives and decision-making processes. So, everyone who uses these tools carries a big responsibility. Being aware of this responsibility makes a real difference compared to the AI approach I mentioned earlier. Even though you teach us that ethics can be taught to machines, for now, human awareness is key.

So, this responsibility means not just creating a tool but doing it ethically, which we believe is very important. Ethical, social, and environmental implications add value to our work—they make it less mechanical, less algorithmic.

This ties back to the concept of harmony, expanding the vision of those who create these tools. We think that's very important.

So yes, we like to talk about explicit awareness—not just about ethics, but as part of humanizing our work.

## Evaluating Utility / Usability

**Q8: Should you label your software product using the ethics label we are proposing, which cards and characteristics would you consider displaying to the user?**

**Q8.1: Why?**

Yes. Among the cards and characteristics in the proposed ethics label, the ones I find most useful and aligned with our work are those related to **transparency**, **explainability**, **data governance**, **privacy**, **accountability**, and **beneficial outcomes**. These are particularly relevant in our context—banking, insurance, and wealth management—where **trust** and **client protection** are absolutely central.

Our users, and their clients, rely on the software we build to make decisions that directly impact their personal financial well-being. These are often based on complex and partially automated information, so it's crucial that we provide tools that help them understand why certain recommendations are made, and how their data is being handled. It's also key to clarify **who is accountable** when something goes wrong and to be able to **trace and remediate** errors transparently.

When it comes to **energy consumption**, this is definitely a concern for us—both from an environmental and a cost perspective. Since most of our solutions are now deployed on the cloud, we don't directly measure energy use. We would love to include this kind of information in the label, but we'd need more support and transparency from cloud providers to do so. Right now, we operate based on common sense—for example, turning off non-essential environments overnight to reduce cost and energy—but we don't have access to objective metrics that would allow us to report this properly.

Regarding the card about **managing adverse effects**, especially in the context of generative AI, we're very aware of the risks—like hallucinations or biased outcomes—and we always include a **validation phase** in any decision-making process supported by AI. We combine this with **explainability**, whenever possible. At the end of the day, it's always the human user who makes the final decision, not the machine.

We also implement mechanisms to reflect **ethical user preferences**—for instance, if a client doesn't want to invest in certain industries for ethical reasons, we make sure those constraints are respected and properly propagated across the system. These choices are explained clearly during the advisory process.

There are also more subtle ethical aspects we take into account, such as potential conflicts of interest—where advisors might earn different fees depending on the financial product they recommend. We believe it's essential to have systems that

**ensure transparency and fairness**, both to protect the client and to support ethical professional practices.

So yes, if we were to use this ethics label, we would definitely consider displaying these characteristics. They not only align with what we already do, but also reinforce **trust**—both with our institutional customers and with their end users. And for us, trust is everything.

### **Anything to add?**

I'm genuinely interested in this topic, also because it's a very current issue that directly affects us and that we'd like to continue exploring. We're not interested in jumping on trends just for the sake of it — we prefer to identify what's truly valuable, especially when it can bring real benefits to our clients. In that sense, this label is definitely welcome. I see it as a useful tool from a developer's perspective as well, because it helps raise awareness and encourages transparency across multiple dimensions of the product. I fully agree with that vision.