

Introductory questions

Q1: How would you describe your company in terms of its mission and activities?

Interviewee_A and Interviewee_B:

We offer consulting services, and we're currently working on developing some software products. Our core business is system integration of market solutions like Salesforce and SAP, which we do together with our parent company.

We were recently acquired by Seidor, a large group with nearly a billion in annual revenue, 8.000 employees, and presence in 44 countries. So, our scope has broadened significantly.

Still, fundamentally, we remain system integrators, providing services. On the Italian side of the company, we're increasingly moving toward generative AI, agents, and similar innovations.

As for locations, we operate not only in L'Aquila, but also had a presence in Rome until recently. We definitely have operational teams in Milan and the province of Teramo.

Q2: How would you describe your professional profile and background?

interviewee_B:

Well, I have a very academic background. I worked as a professor and reached the level of Assistant Professor in Cuba, where I spent around five years at the university.

I was always involved in experimental development and research projects, primarily within the academic sphere and, to a limited extent, outside of it.

Then I did a four-year PhD in Computer Science, and for the past year and a half, I've been working at Gunpowder, primarily as a developer, but I also coordinate some activities.

interviewee_A:

He's also the supervisor of all the research activities, I should add.

As for me, well, we both forgot to mention that we're originally electronics engineers—that's our main degree.

After university, I went straight into the job market: I spent 14 years in the telecom sector with Vodafone in Milan, followed by about 10 years in Modena in luxury automotive manufacturing, specifically with Maserati and Lamborghini.

Since 2021, I've been with Gunpowder. Technically, my contract says I'm the Chief Operating Officer, but in reality, I'm currently focused on managing and developing the business unit for manufacturing and automotive.

We structure ourselves around market-oriented business units and cross-cutting competence centers. These centers provide the knowledge we need to get things done.

In terms of knowledge areas, as I mentioned earlier, I'm particularly involved in IoT, Industry 4.0, and Industry 5.0, with a strong personal interest in leading and supporting interviewee_B's group as well.

Q3: Can you briefly describe the software product your company develops, and the role of AI within it?

Interviewee_B:

We do have a software product, although it's still in beta. As said, we're not a typical software house — we're more of a consulting firm. We identified a market niche for a product that enables the use of generative AI, offering an experience similar to ChatGPT, but deployable also in **on-premise** environments. This makes it suitable for public administrations or companies that don't want their data to ever leave their infrastructure, yet still want to leverage the advantages of generative AI — even if it means using less powerful models.

Our solution is primarily an **aggregator of AI services**. It provides a unified interface to connect with different AI providers or services — including LLMs, predictive models, embedding models, etc. — whether these are deployed on-premise or externally. So, our platform can work both with and without an internet connection.

The role of AI is **foundational** to this product. The solution is not tied to any specific provider — it's provider-agnostic — and emphasizes local deployment. Besides aggregating providers or offering self-hosted models, it allows users to access multiple services from a single interface. For example, users can generate documents in specific formats, like with a company's letterhead, or produce receipts with the help of generative components.

The product also includes features for managing shared knowledge bases, with permissions and roles.

AI, for us, is the foundation. What we do is **treat AI as an abstract foundation** — a layer — on top of which we build custom solutions, without locking into a specific vendor like OpenAI or Google.

Interviewee_A:

Yes, I agree. Even though our platform is still in beta, we can definitely use it as a use case for the rest of the interview.

We basically **wrap multiple AI providers** — like OpenAI, Gemini, Anthropic — and we aggregate their capabilities into our interface. The actual intelligence is in the orchestration layer we built, and that's where we make our design choices and add value.

That's the part we control and develop — the layer that integrates everything and offers a consistent user experience. So yes, we can certainly refer to this product and we can definitely use it as a use case.

Q4: Have you ever heard of privacy or nutrition labels?

interviewee_A: I have never seen the privacy label before.

interviewee_B: Me neither.

Evaluating Comprehensibility / Understandability

Q5: Is the proposed ethics label easily understandable?

Interviewee_A:

In my opinion, if you want my take — and I have the Ethics label, the one in table format, right in front of me — I glanced at it and maybe it's a bit too much...It's kind of... it should be a bit more concise if someone wants to sell it, it should be shortened a bit.

However, then the ethics label as in the application example is clear, concise and not verbose.

Interviewee_B:

Yes, yes, it is easy to understand. The layout is clear, and the content is straightforward.

Evaluating Appropriateness

Q6: Does the proposed label appropriately address Human, Societal, and Environmental drivers?

interviewee_B:

Yes, I would say so. Since that's where you started from, it clearly had an important influence, which is indeed reflected in the Ethics label.

interviewee_A:

I got a bit lost here. When we talk about Human, Societal, and Environmental drivers, are we including things like cultural values or historical heritage?

So my point is: are we also addressing that kind of dimension here?

Later I noticed the label includes something like "*cultural sensitiveness*" under the Societal column. But I don't think that's explicitly present in the list of drivers, which may be a bit of a gap. So, I'd say the label mostly covers the drivers—perhaps not entirely one-to-one—but overall, yes. Some aspects like cultural issues may be underrepresented.

Still, I'd say it's more yes than no.

Evaluating Usefulness

Q7: From a developer's perspective, is the proposed ethics label useful for expressing the ethical quality of your product?

Q7.1: Why?

Interviewee_B:

If all the fields of the label are not mandatory, I have concerns because if I put the label as a developer, or better as a company, the ethics of the label is tied to the ethics of the company. So, it is not the ethics of a certifying body or the consumer, but the ethics of the company. What it reflects is the mirror of the company's ethics.

If all elements are not mandatory, I can just put the ones I think are more compliant. But as a consumer, what interests me more is to look at those that are missing, because they are not there.

A generalization of this label could foresee a small note explaining why the missing elements are not there. This would give more completeness, not because it improves self-certified ethics, but because it forces the company to make a declaration about all elements that can be evaluated.

There is also the issue of ethics washing, which this kind of label tries to overcome.

If I, as self-certified, must express a judgment on all points, that definitely improves transparency about what is being declared.

Maybe to popularize the use of such a label, one could think of how this label can be built by different parties. For example, in open source products, even if large companies are involved, many things can be at least estimated, if not all. Maybe through sub-certifications, others could be estimated.

One use case after defining the label could be a distributed certification mechanism, so it is not self-certified, because the problem with self-certification is that it reflects the ethics of whoever writes it—in this case, the company.

Ultimately, consumers judge the company's ethics over time as issues emerge.

I also understand that not all aspects can be measured with exact metrics.

Interviewee_A:

I don't have much to add here.

Evaluating Utility / Usability

Q8: Should you label your software product using the ethics label we are proposing, which cards and characteristics would you consider displaying to the user?

Q8.1: Why?

Interviewee_B:

I believe it's important to show all characteristics—both the positive and the negative ones—because the overall value for the user comes only when everything is declared. If certain characteristics don't apply, I would still include them and specify why they are not applicable. For example, *Water Consumption* doesn't apply directly to our product, and maybe only *Energy Consumption* is relevant from the *Sustainability* category.

As for *Greenhouse Gas Emissions*, there might be indirect effects, but I would still include it and explain the limitations. When it comes to *Human and Societal* aspects, I think we can reflect on all of them. *Transparency and Explainability* might be harder due to the use of LLMs, but I would still address it and acknowledge the limits.

The important point is to take a stance on every aspect—even when it's subjective or hard to measure—because it shows a commitment. Leaving something out reduces the label's reliability.

Regarding *Inclusive and Participatory Design*, I think our tool does contribute to that. Even *Society and Democracy* might apply in a broader sense, because bringing generative AI into areas where current tools like ChatGPT aren't usable helps democratize access and automate tasks.

It might also be helpful to have **guidelines**—not necessarily metrics—for each subcategory to support clearer evaluation.

Interviewee_A:

I agree with what was said—we should include both positive and negative characteristics, even those that don't apply. For example, *Water Consumption* clearly doesn't apply to us.

In general, I think most of the *Human and Societal* cards would apply to our product, and I'd include them. Even if some aspects are harder to evaluate, I'd still mention them and explain the context.

Anything to add?

Interviewee_B:

No, I think I've said pretty much everything. I would emphasize trying to require that a judgment be expressed on every item, and wherever possible, to automate the evaluation—at least for open-source products—as a potential use case.

From a consumer's perspective, that's what I would value the most: having a full picture of the product, even those aspects that might be hard to disclose when trying to sell it.

It's similar to food labeling—if your product has a lot of sugar, you're probably not thrilled to say it, but you have to. That's the point.

Interviewee_A:

The only thing I'd add is something I mentioned earlier—understanding the cultural aspect. You might ask what that has to do with ethics, but I believe it matters. Thank you again, and congratulations—your work is genuinely interesting. I hope it gains traction and goes far.